

In [50]:

```
1 import pandas as pd
2 import numpy as np
3 import matplotlib.pyplot as plt
4 import seaborn as sns
5 from sklearn.model_selection import train_test_split,GridSearchCV,RandomizedSearchCV
6 from sklearn.tree import DecisionTreeRegressor
7 from sklearn.metrics import mean_squared_error,mean_absolute_error,r2_score
8 import statsmodels.api as sm
9 from scipy.stats import shapiro, kstest,normaltest
10 import pickle
11 import json
12 import os
```

Step 1: Problem statement:

1	To predict the organisational level co2 emmission based on the fule consumption vs co2 emission data, collected for different vehicles at different traffic conditions.
2	
3	Understanding the Data
4	Model 4WD/4X4 = Four-wheel drive
5	AWD = All-wheel drive
6	FFV = Flexible-fuel vehicle
7	SWB = Short wheelbase
8	LWB = Long wheelbase
9	EWB = Extended wheelbase
10	Transmission A = automatic
11	AM = automated manual
12	AS = automatic with select shift
13	AV = continuously variable
14	M = manual
15	3 - 10 = Number of gears
16	Fuel type X = regular gasoline
17	Z = premium gasoline
18	D = diesel
19	E = ethanol (E85)
20	N = natural gas
21	Fuel consumption City and highway fuel consumption ratings are shown in
22	litres per 100 kilometres (L/100 km) - the combined rating (55% city, 45% hwy)
23	is shown in L/100 km and in miles per imperial gallon (mpg)
24	CO2 emissions the tailpipe emissions of carbon dioxide (in grams per kilometre)
25	for combined city and highway driving
26	

1	Step 2: Data Gathering.
2	data gathering id a task data engineer,
3	the current data set is downloaded from kaggle

Step 2: Data Gathering

```
In [2]: 1 df=pd.read_csv('CO2 Emissions_Canada.csv')
        2 df
```

Out[2]:

	Make	Model	Vehicle Class	Engine Size(L)	Cylinders	Transmission	Fuel Type	Fuel Consumption City (L/100 km)	Fuel Consumption Hwy (L/100 km)	Cor Co
0	ACURA	ILX	COMPACT	2.0	4	AS5	Z	9.9	6.7	
1	ACURA	ILX	COMPACT	2.4	4	M6	Z	11.2	7.7	
2	ACURA	ILX HYBRID	COMPACT	1.5	4	AV7	Z	6.0	5.8	
3	ACURA	MDX 4WD	SUV - SMALL	3.5	6	AS6	Z	12.7	9.1	
4	ACURA	RDX AWD	SUV - SMALL	3.5	6	AS6	Z	12.1	8.7	
...	
7380	VOLVO	XC40 T5 AWD	SUV - SMALL	2.0	4	AS8	Z	10.7	7.7	
7381	VOLVO	XC60 T5 AWD	SUV - SMALL	2.0	4	AS8	Z	11.2	8.3	
7382	VOLVO	XC60 T6 AWD	SUV - SMALL	2.0	4	AS8	Z	11.7	8.6	
7383	VOLVO	XC90 T5 AWD	SUV - STANDARD	2.0	4	AS8	Z	11.2	8.3	
7384	VOLVO	XC90 T6 AWD	SUV - STANDARD	2.0	4	AS8	Z	12.2	8.7	

7385 rows × 12 columns

Step 3: EDA & Feature Engineering

```
In [3]: 1 df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7385 entries, 0 to 7384
Data columns (total 12 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Make                                  7385 non-null   object
1   Model                                7385 non-null   object
2   Vehicle Class                        7385 non-null   object
3   Engine Size(L)                       7385 non-null   float64
4   Cylinders                            7385 non-null   int64
5   Transmission                         7385 non-null   object
6   Fuel Type                            7385 non-null   object
7   Fuel Consumption City (L/100 km)     7385 non-null   float64
8   Fuel Consumption Hwy (L/100 km)      7385 non-null   float64
9   Fuel Consumption Comb (L/100 km)     7385 non-null   float64
10  Fuel Consumption Comb (mpg)           7385 non-null   int64
11  CO2 Emissions(g/km)                  7385 non-null   int64
dtypes: float64(4), int64(3), object(5)
memory usage: 692.5+ KB
```

In [4]:

1

x=df.drop(['Make','Model','CO2 Emissions(g/km)'], axis=1)

2

x

Out[4]:

	Vehicle Class	Engine Size(L)	Cylinders	Transmission	Fuel Type	Fuel Consumption City (L/100 km)	Fuel Consumption Hwy (L/100 km)	Fuel Consumption Comb (L/100 km)	Consum
0	COMPACT	2.0	4	AS5	Z	9.9	6.7	8.5	
1	COMPACT	2.4	4	M6	Z	11.2	7.7	9.6	
2	COMPACT	1.5	4	AV7	Z	6.0	5.8	5.9	
3	SUV - SMALL	3.5	6	AS6	Z	12.7	9.1	11.1	
4	SUV - SMALL	3.5	6	AS6	Z	12.1	8.7	10.6	
...	
7380	SUV - SMALL	2.0	4	AS8	Z	10.7	7.7	9.4	
7381	SUV - SMALL	2.0	4	AS8	Z	11.2	8.3	9.9	
7382	SUV - SMALL	2.0	4	AS8	Z	11.7	8.6	10.3	
7383	SUV - STANDARD	2.0	4	AS8	Z	11.2	8.3	9.9	
7384	SUV - STANDARD	2.0	4	AS8	Z	12.2	8.7	10.7	

7385 rows × 9 columns

In [5]:

1

y=df[['CO2 Emissions(g/km)']]

2

y

Out[5]:

	CO2 Emissions(g/km)
0	196
1	221
2	136
3	255
4	244
...	...
7380	219
7381	232
7382	240
7383	232
7384	248

7385 rows × 1 columns

```
In [6]: 1 x.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7385 entries, 0 to 7384
Data columns (total 9 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Vehicle Class                        7385 non-null   object
1   Engine Size(L)                      7385 non-null   float64
2   Cylinders                           7385 non-null   int64
3   Transmission                        7385 non-null   object
4   Fuel Type                           7385 non-null   object
5   Fuel Consumption City (L/100 km)    7385 non-null   float64
6   Fuel Consumption Hwy (L/100 km)     7385 non-null   float64
7   Fuel Consumption Comb (L/100 km)    7385 non-null   float64
8   Fuel Consumption Comb (mpg)         7385 non-null   int64
dtypes: float64(4), int64(2), object(3)
memory usage: 519.4+ KB
```

```
In [7]: 1 x= pd.get_dummies(x, columns=['Vehicle Class'])
```

```
In [8]: 1 x = pd.get_dummies(x, columns=['Transmission'])
```

```
In [9]: 1 x
```

Out[9]:

	Engine Size(L)	Cylinders	Fuel Type	Fuel Consumption City (L/100 km)	Fuel Consumption Hwy (L/100 km)	Fuel Consumption Comb (L/100 km)	Fuel Consumption Comb (mpg)	Vehicle Class_COMPACT	C
0	2.0	4	Z	9.9	6.7	8.5	33	True	
1	2.4	4	Z	11.2	7.7	9.6	29	True	
2	1.5	4	Z	6.0	5.8	5.9	48	True	
3	3.5	6	Z	12.7	9.1	11.1	25	False	
4	3.5	6	Z	12.1	8.7	10.6	27	False	
...	
7380	2.0	4	Z	10.7	7.7	9.4	30	False	
7381	2.0	4	Z	11.2	8.3	9.9	29	False	
7382	2.0	4	Z	11.7	8.6	10.3	27	False	
7383	2.0	4	Z	11.2	8.3	9.9	29	False	
7384	2.0	4	Z	12.2	8.7	10.7	26	False	

7385 rows × 50 columns

```
In [10]: 1 x.replace({True:1,False:0},inplace=True)
        2 x
```

Out[10]:

	Engine Size(L)	Cylinders	Fuel Type	Fuel Consumption City (L/100 km)	Fuel Consumption Hwy (L/100 km)	Fuel Consumption Comb (L/100 km)	Fuel Consumption Comb (mpg)	Vehicle Class_COMPACT	C
0	2.0	4	Z	9.9	6.7	8.5	33	1	
1	2.4	4	Z	11.2	7.7	9.6	29	1	
2	1.5	4	Z	6.0	5.8	5.9	48	1	
3	3.5	6	Z	12.7	9.1	11.1	25	0	
4	3.5	6	Z	12.1	8.7	10.6	27	0	
...
7380	2.0	4	Z	10.7	7.7	9.4	30	0	
7381	2.0	4	Z	11.2	8.3	9.9	29	0	
7382	2.0	4	Z	11.7	8.6	10.3	27	0	
7383	2.0	4	Z	11.2	8.3	9.9	29	0	
7384	2.0	4	Z	12.2	8.7	10.7	26	0	

7385 rows × 50 columns

```
In [11]: 1 x['Fuel Type'].unique()
```

Out[11]: array(['Z', 'D', 'X', 'E', 'N'], dtype=object)

```
In [12]: 1 x['Fuel Type'].replace({'Z':3, 'D':5, 'X':4, 'E':2, 'N':1},inplace=True)
```

```
In [13]: 1 x.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7385 entries, 0 to 7384
Data columns (total 50 columns):
#   Column                                                                 Non-Null Count  Dtype
---  -
0   Engine Size(L)                                                         7385 non-null   float64
1   Cylinders                                                             7385 non-null   int64
2   Fuel Type                                                             7385 non-null   int64
3   Fuel Consumption City (L/100 km)                                     7385 non-null   float64
4   Fuel Consumption Hwy (L/100 km)                                     7385 non-null   float64
5   Fuel Consumption Comb (L/100 km)                                    7385 non-null   float64
6   Fuel Consumption Comb (mpg)                                          7385 non-null   int64
7   Vehicle Class_COMPACT                                                7385 non-null   int64
8   Vehicle Class_FULL-SIZE                                              7385 non-null   int64
9   Vehicle Class_MID-SIZE                                               7385 non-null   int64
10  Vehicle Class_MINICOMPACT                                             7385 non-null   int64
11  Vehicle Class_MINIVAN                                                7385 non-null   int64
12  Vehicle Class_PICKUP TRUCK - SMALL                                  7385 non-null   int64
13  Vehicle Class_PICKUP TRUCK - STANDARD                              7385 non-null   int64
14  Vehicle Class_SPECIAL PURPOSE VEHICLE                              7385 non-null   int64
15  Vehicle Class_STATION WAGON - MID-SIZE                             7385 non-null   int64
16  Vehicle Class_STATION WAGON - SMALL                                7385 non-null   int64
17  Vehicle Class_SUBCOMPACT                                             7385 non-null   int64
18  Vehicle Class_SUV - SMALL                                           7385 non-null   int64
19  Vehicle Class_SUV - STANDARD                                        7385 non-null   int64
20  Vehicle Class_TWO-SEATER                                             7385 non-null   int64
21  Vehicle Class_VAN - CARGO                                            7385 non-null   int64
22  Vehicle Class_VAN - PASSENGER                                       7385 non-null   int64
23  Transmission_A10                                                     7385 non-null   int64
24  Transmission_A4                                                       7385 non-null   int64
25  Transmission_A5                                                       7385 non-null   int64
26  Transmission_A6                                                       7385 non-null   int64
27  Transmission_A7                                                       7385 non-null   int64
28  Transmission_A8                                                       7385 non-null   int64
29  Transmission_A9                                                       7385 non-null   int64
30  Transmission_AM5                                                      7385 non-null   int64
31  Transmission_AM6                                                      7385 non-null   int64
32  Transmission_AM7                                                      7385 non-null   int64
33  Transmission_AM8                                                      7385 non-null   int64
34  Transmission_AM9                                                      7385 non-null   int64
35  Transmission_AS10                                                     7385 non-null   int64
36  Transmission_AS4                                                      7385 non-null   int64
37  Transmission_AS5                                                      7385 non-null   int64
38  Transmission_AS6                                                      7385 non-null   int64
39  Transmission_AS7                                                      7385 non-null   int64
40  Transmission_AS8                                                      7385 non-null   int64
41  Transmission_AS9                                                      7385 non-null   int64
42  Transmission_AV                                                       7385 non-null   int64
43  Transmission_AV10                                                     7385 non-null   int64
44  Transmission_AV6                                                      7385 non-null   int64
45  Transmission_AV7                                                      7385 non-null   int64
46  Transmission_AV8                                                      7385 non-null   int64
47  Transmission_M5                                                       7385 non-null   int64
48  Transmission_M6                                                       7385 non-null   int64
49  Transmission_M7                                                       7385 non-null   int64
dtypes: float64(4), int64(46)
memory usage: 2.8 MB
```

There are no any null/missing values so missing values handling is not required

Here we are going to use Decision Tree Model, so Outlier Handling / Scaling is not required

Step 4: Feature Selection

there are no any assumptions over data so we will be performing this step after model training and eveluation

step 5 Model Training

```
In [15]: 1 x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.2, random_state=1)
        2 x_train
```

Out[15]:

	Engine Size(L)	Cylinders	Fuel Type	Fuel Consumption City (L/100 km)	Fuel Consumption Hwy (L/100 km)	Fuel Consumption Comb (L/100 km)	Fuel Consumption Comb (mpg)	Vehicle Class_COMPACT	C
580	2.4	4	4	11.1	8.3	9.8	29	0	
3998	2.0	4	4	11.2	7.6	9.8	29	0	
2228	2.0	4	3	11.2	8.4	9.9	29	0	
2954	2.5	4	4	8.7	6.5	7.7	37	0	
4	3.5	6	3	12.1	8.7	10.6	27	0	
...
905	3.4	6	3	11.9	8.6	10.4	27	0	
5192	2.0	4	3	9.3	7.3	8.4	34	0	
3980	2.0	4	3	10.7	8.5	9.7	29	0	
235	2.4	4	4	12.2	8.6	10.6	27	0	
5157	3.0	6	3	11.8	8.7	10.4	27	0	

5908 rows × 50 columns

```
In [18]: 1 dt_reg=DecisionTreeRegressor()
        2 dt_reg
```

Out[18]: DecisionTreeRegressor()

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [19]: 1 dt_reg.fit(x_train,y_train)
```

Out[19]: DecisionTreeRegressor()

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [21]: 1 y_pred=dt_reg.predict(x_test)
        2 y_pred
```

Out[21]: array([202., 214., 174., ..., 177., 115., 194.])

```
In [22]: 1 y_pred_train = dt_reg.predict(x_train)
        2 y_pred_train
```

Out[22]: array([225., 230., 232., ..., 226., 244., 246.])

Step 6 Model Evaluation: model is being evaluated in this step

```
In [27]: 1 mse_test = mean_squared_error(y_test,y_pred)
        2 mse_test
```

Out[27]: 7.408273443834266

```
In [28]: 1 mae_test = mean_absolute_error(y_test,y_pred)
        2 mae_test
```

Out[28]: 1.5294978022804697

In [29]:

1 np.sqrt(mse_test)

Out[29]:

2.721814366159872

In [30]:

1 mse_train = mean_squared_error(y_train,y_pred_train)
2 mse_train

Out[30]:

1.2366056651728623

In [31]:

1 mae_train = mean_absolute_error(y_train,y_pred_train)
2 mae_train

Out[31]:

0.42930610632878735

In [32]:

1 np.sqrt(mse_train)

Out[32]:

1.1120277268003989

In [33]:

1 r2score_test=r2_score(y_test,y_pred)
2 r2score_test

Out[33]:

0.9978242028977429

In [34]:

1 r2score_train=r2_score(y_train,y_pred_train)
2 r2score_train

Out[34]:

0.9996392293300826

1 conclusion: model is performing well as
2 1. we have low bias and low variance
3 2. R2_score is > 0.99
4 3. MSE and MAE values are low for both training and testing data.

Hyper parameter tuning

In [51]:

1 df_reg_h = DecisionTreeRegressor()
2 param_grid = {'criterion' : ['squared_error', 'absolute_error'],
3 'max_depth' : np.arange(5,15),
4 'min_samples_split' : np.arange(3,8),
5 'min_samples_leaf' : np.arange(2,5)}
6 gscv = RandomizedSearchCV(df_reg_h,param_grid,cv=5,n_jobs=-1)

In [52]:

1 gscv.fit(x_train,y_train)

Out[52]:

RandomizedSearchCV

RandomizedSearchCV(cv=5, estimator=DecisionTreeRegressor(), n_jobs=-1,
param_distributions={'criterion': ['squared_error',
'absolute_error'],
'max_depth': array([5, 6, 7, 8, 9, 10,
11, 12, 13, 14]),
'min_samples_leaf': array([2, 3, 4]),
'min_samples_split': array([3, 4, 5, 6,
7])})

estimator: DecisionTreeRegressor
DecisionTreeRegressor()
DecisionTreeRegressor()

In [53]:

1 gscv.best_estimator_

Out[53]:

DecisionTreeRegressor
DecisionTreeRegressor(max_depth=9, min_samples_leaf=2, min_samples_split=3)

In [54]:

```
1 dt_reg_hyp = DecisionTreeRegressor(max_depth=9, min_samples_leaf=2, min_samples_spli
2 dt_reg_hyp.fit(x_train,y_train)
```

Out[54]:

DecisionTreeRegressor

DecisionTreeRegressor(max_depth=9, min_samples_leaf=2, min_samples_split=3)

In [55]:

```
1 y_pred=dt_reg_hyp.predict(x_test)
2 y_pred_train = dt_reg_hyp.predict(x_train)
3 mse_test = mean_squared_error(y_test,y_pred)
4 print('mse_test', mse_test)
5 mae_test = mean_absolute_error(y_test,y_pred)
6 print('mae_test', mae_test)
7 print('RMSE Test', np.sqrt(mse_test))
8 r2score_test=r2_score(y_test,y_pred)
9 print('r2score_test',r2score_test)
10 mse_train = mean_squared_error(y_train,y_pred_train)
11 print('mse_train', mse_train)
12 mae_train = mean_absolute_error(y_train,y_pred_train)
13 print('mae_train', mae_train)
14 print('RMSE Train', np.sqrt(mse_train))
15 r2score_train=r2_score(y_train,y_pred_train)
16 print('r2score_train',r2score_train)
```

mse_test 7.369223276074442

mae_test 2.064009165325812

RMSE Test 2.71463133336268

r2score_test 0.9978356718645

mse_train 6.131035844703403

mae_train 1.8566362075012992

RMSE Train 2.476092858659263

r2score_train 0.9982113150770079

```
1 conclusion: By using best parameters from hyper parameter tuning,
2 model is performing well as
3 1. we have low bias and low variance
4 2. R2_score is > 0.99
5 3. MSE and MAE values are low for both training and testing data.
```

Creating get data function, pickle and json data file

In [35]:

```
1 def get_input_row(make,model,Vehicle_Class, Engine_Size, Cylinders,Transmission, Fuel
2 ,Fuel_Consumption_City1, Fuel_Consumption_Hwy1
3 ,Fuel_Consumption_Comb2, Fuel_Consumption_Comb3):
4 df1=pd.DataFrame(np.zeros(shape=(50)))
5 df1.index=x.columns
6 df2=df1.T
7 df2['Engine Size(L)']=Engine_Size
8 df2['Cylinders']=Cylinders
9 df2['Fuel Consumption City (L/100 km)']=Fuel_Consumption_City1
10 df2['Fuel Consumption Hwy (L/100 km)']=Fuel_Consumption_Hwy1
11 df2['Fuel Consumption Comb (L/100 km)']=Fuel_Consumption_Comb2
12 df2['Fuel Consumption Comb (mpg)']=Fuel_Consumption_Comb3
13 df2['Fuel Type']=Fuel_Type
14 col_name='Vehicle Class_'+ Vehicle_Class
15 df2[col_name]= 1
16 col_name1='Transmission_'+ Transmission
17 df2[col_name1]=1
18 df2['Fuel Type'].replace({'Z':3, 'D':5, 'X':4, 'E':2, 'N':1},inplace=True)
19 return df2
```

In [56]:

```
1 input_df=get_input_row('Suraj', 'SUV', 'COMPACT', 2.0, 4, 'AS5', 'Z', 9.9, 6.7, 8.6,
2 y_predicted = dt_reg_hyp.predict(input_df)
3 predicted_co2_emmission = y_predicted[0]
```

In [57]:

```
1 with open('DT_regression.pkl','wb') as f:
2 pickle.dump(dt_reg_hyp,f)
```

In [39]:

```
1 dict1 = {'columns_x' : x.columns.to_list() }  
2 with open('project_data.json','w') as f:  
3     json.dump(dict1,f)
```